

---

# COEVOLUTION NEUROEVOLUTION

---

Luca Manzoni

---

---

# COEVOLUTION

---



---

# COEVOLUTION: IDEAS

---

- The fitness of a single individual is now influenced by “external factors”:
    - Performance against the other individual of the population
    - “Collective performance”: the entire population counts
    - By the similarity to other individuals: too many similar individuals are “bad”.
-



---

# THE ROLE OF FITNESS

---

- Usually we are interested in maximising/minimising a certain fitness value...
  - ...but now the fitness is influenced by the other individuals
  - **Absolute fitness:** the fitness that we actually want to optimise
  - **Relative fitness:** the fitness that the algorithm is optimising, hopefully also improving the absolute fitness
-



---

# ONE-POPULATION COMPETITIVE COEVOLUTION

---

- Usually employed when producing individuals that play games
  - Works when evolving agents where the quality of the solution can be assessed by making them play one against the other
  - Here evaluating the “real fitness” might be difficult
-




# WHY?

Fitness evaluation  
with a “competent adversary”

Population at generation 0

Fitness  
Number of victories






# WHY?

Fitness evaluation  
with a “not-too-competent adversary”

Population at generation 100

Fitness  
Number of victories





---

# PROBLEM OF USING A FIXED OPPONENT

---

- Evaluating the fitness of the individuals high be difficult if the individuals are too weak or too strong w.r.t. the opponent
  - This might be solved by using a collection of fixed opponents...
  - ...or by using directly other individuals as opponents
-



---

# INTERNAL VS EXTERNAL FITNESS

---

Individuals might get better  
at winning against  
other individuals...

 **Internal fitness**

**External fitness**

...but not against real opponents



**Progress should be evaluated with both fitnesses**


---





---

# EVALUATION: PAIRWISE

---



Only make individual  $n$  compete with individual  $n+1$



**+only  $O(n)$  tests**

**-very dependant on the adversary**


---




---

# EVALUATION: COMPLETE

---



Every individual competes with every other individual



**+precise assessment of the fitness**

**-with  $O(n^2)$  tests**


---





---

# EVALUATION: K-FOLD

---



Every individual competes with k randomly selected individual



**+tunable number of tests**

**-some individuals might be tested more than others**


---




---

# EVALUATION: SINGLE-ELIMINATION TOURNAMENT

---



Every individual proceeds with the competitions until it is beaten



**+Only  $O(n)$  tests**

**+Better individuals are tested more times**

**-Bad pairing might penalise some individuals**


---




---

# FITLESSNESS SELECTION

---



We never compute the fitness:  
in tournament selection the selected individual is  
the winner of the tournament





---


# TWO-POPULATIONS COMPETITIVE COEVOLUTION

---

- Sometimes we want two populations to compete against each other to improve together
  - **Primary population:** the individuals we are actually interested in
  - **Alternative (foil) population:** individuals that try to beat/foil the individuals in the primary population
-



# SORTING NETWORKS




A comparator exchanges the two values if the one at the top is bigger than the one at the bottom.


*The optimal depth of a sorting network for  $n$  elements is known only for small  $n$*



# EVOLVING OF SORTING NETWORKS



Population  
of sorting networks



Population  
of "difficult to sort" arrays



---

# HOW TO DEFINE THE FITNESS

---

## **For sorting networks**

Number of correctly  
sorted arrays



Networks that sort well

## **For arrays**

Number of “foiled”  
networks



Arrays that represents  
corner cases for networks


The two fitnesses are “in conflict”,  
thus the coevolution is competitive

---



# HOW THE EVOLUTION IS PERFORMED (I)

P and Q are always changing,  
how to assess the fitness of the individuals?



Shuffle and pairing

Repeat k times


Avoid performing the  
same evaluation more  
than once



# HOW THE EVOLUTION IS PERFORMED (II)

To make evolution easier we can use the same procedure but:

- 1)  $Q$  at time  $t$  is compared with  $P$  at time  $t-1$
- 2)  $P$  at time  $t$  is compared with  $Q$  at time  $t-1$





---

# HOW THE EVOLUTION IS PERFORMED (III)

---

- The comparison with the previous generation allows to perform some further tuning:
  - We can compare each individual of  $P$  (resp.,  $Q$ ) at time  $t$  with the  $k$  best individuals in population  $Q$  (resp.,  $P$ ) at time  $t-1$
  - We can combine the two approaches and select  $k_1$  individuals from the best and  $k_2$  randomly
-



---

# LOSS OF GRADIENT

---

- Two-populations competitive coevolution is not without problems:
  - If one of the two populations is too strong w.r.t. the other then the internal fitness gives no information...
  - ...thus, the evolution process proceeds without any guidance (i.e., the selection is almost uniform among the individuals)
  - It might be useful to stop the evolution of the stronger population and allow the weaker one to “catch up”
-



---

# N-POPULATION COOPERATIVE COEVOLUTION

---

- Sometimes a solution can be decomposed into multiple interacting sub-solutions, each one with its own characteristics
  - Some examples: robot soccer, any multiple player game with more than one role
  - It is always possible to create an enormous individual, but it might be useful to have multiple interacting populations
-



---

# HOW TO ASSESS THE FITNESS?

---

## **Parallel methods**

The evolution happens  
for all populations  
at the same time

The same methods of  
competitive evolution

As usual, we need to specify  
how individuals are paired

## **Sequential methods**

Populations are evolved  
one at a time

This is a case of  
Alternating Optimisation

Not actually suitable  
for competitive evolution

---



---

# POSSIBLE DRAWBACKS

---

Suppose that you want to evolve a soccer team



It is not doing anything useful  
for the global solution

Still has a good fitness because the  
other individuals are strong

---



---

# POSSIBLE DRAWBACKS

---

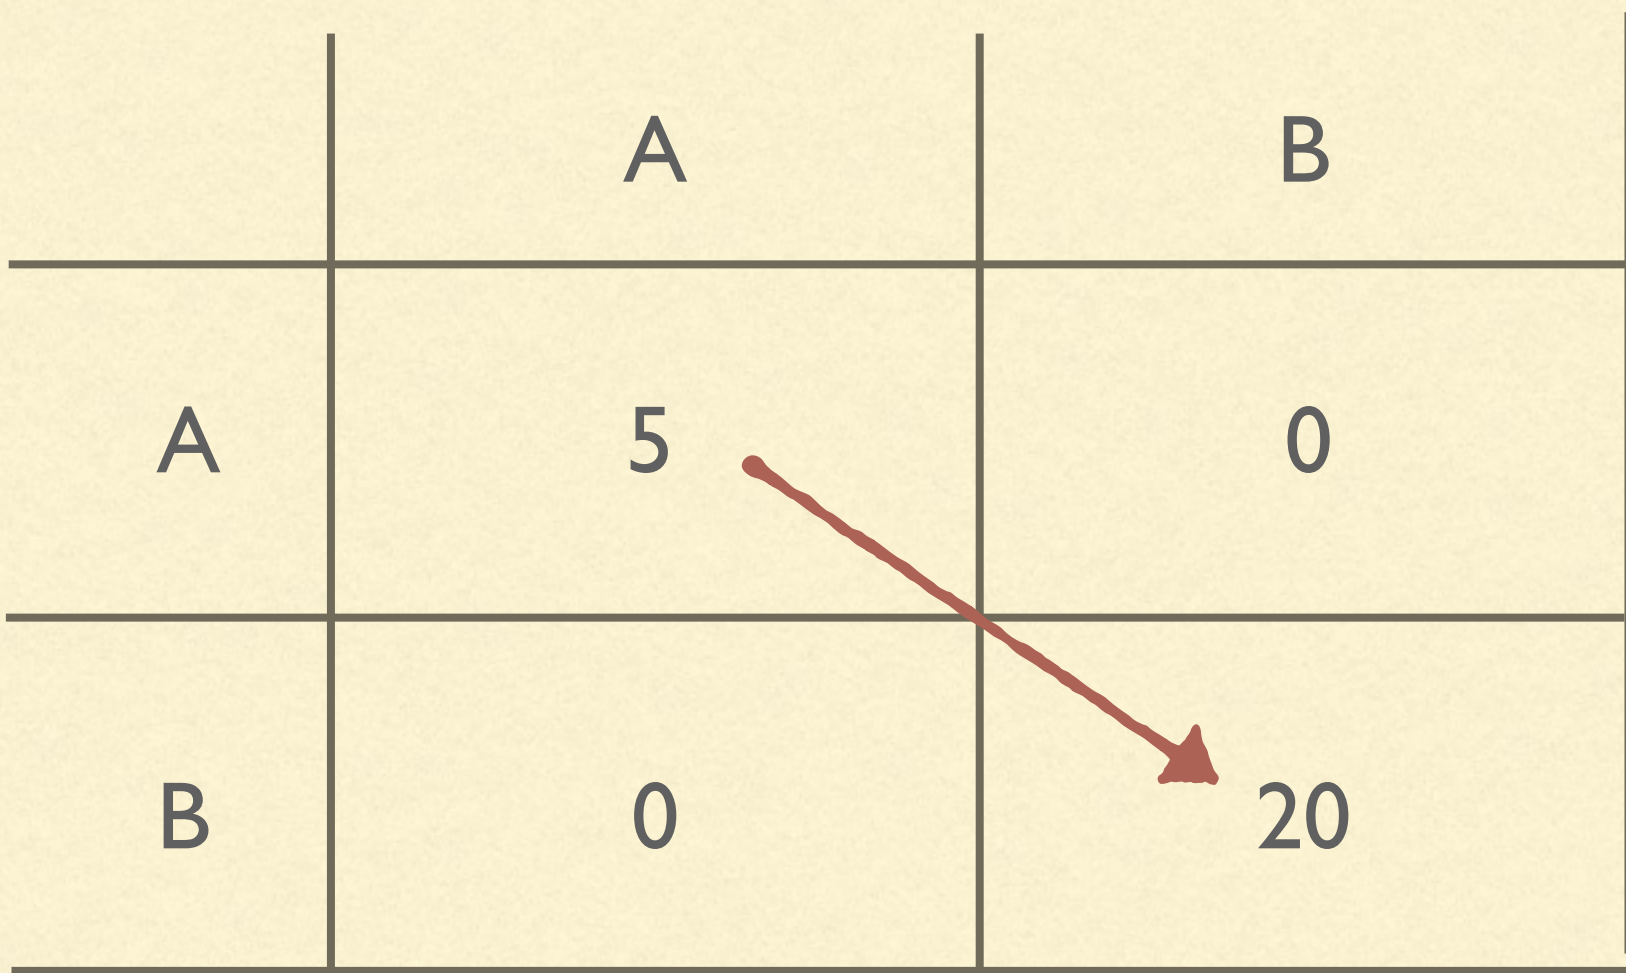
- Overgeneralisation: every individual is decent at everything...
  - ...but good at nothing (but their average performance are good)
  - One possible solution is to compute the fitness as the max (resp., min) among all the  $k$  tests, instead of taking the average, but it is still not a perfect solution
-



# POSSIBLE DRAWBACKS

Behaviours for individuals in population 1

	A	B
A	5	0
B	0	20



The image shows a 2x2 payoff matrix for two populations. The columns represent behaviors A and B for population 1, and the rows represent behaviors A and B for population 2. The payoffs are: (A,A)=5, (A,B)=0, (B,A)=0, (B,B)=20. A red arrow points from the cell (A,A) to the cell (B,B), indicating a transition from a local optimum to a global optimum.


To move from the local optimum 5 the two population must change behaviour “at the same time”



---

# DIVERSITY MAINTENANCE

---






---

# DIVERSITY MAINTENANCE

---






---

# DIVERSITY MAINTENANCE

---

But we would like to preserve some diversity among the solutions





---

# DIVERSITY MAINTENANCE

---

- One of the problems to avoid in optimisation is premature convergence to a sub-optimal solution
  - One possible way to avoid this is to increase the size of the population or “fiddling” with the parameters of the algorithm
  - Otherwise, it is possible to modify the algorithm to enforce some kind of diversity in the population.
-



---

# MEASURING SIMILARITY

---

- **Genotype**: similar construction (e.g., small Hamming distance for sequences of bits)
  - **Phenotype**: similar behaviour (the solution encoded is similar). For example  $(+ \ x \ x)$  and  $(* \ x \ 2)$  have the same phenotype even if they are represented differently
  - **Fitness**: similar fitness value (useful in the case of multi-objective optimisation)
-



# FITNESS SHARING

Two individuals that are too similar have their fitness reduced by having to “share” part of it

Degree of punishment  $\alpha > 0$

Threshold

$$s(x, y) = \begin{cases} 1 - \left( \frac{d(x, y)}{\sigma} \right)^\alpha & \text{if } d(x, y) \leq \sigma \\ 0 & \text{otherwise} \end{cases}$$

Punishment of  $x$  for being too similar to  $y$

The diagram illustrates the fitness sharing formula with several annotations. A green arrow points from the exponent  $\alpha$  in the formula to the text 'Degree of punishment  $\alpha > 0$ '. A red arrow points from the threshold  $\sigma$  in the formula to the text 'Threshold'. A blue arrow points from the function  $s(x, y)$  to the text 'Punishment of  $x$  for being too similar to  $y$ '. The function  $s(x, y)$  is circled in blue, and the threshold  $\sigma$  is circled in red.



---

# FITNESS SHARING

---

The new value of the fitness is computed using the punishment:

Raw fitness

Scaling factor  
 $\beta > 1$

$$f(x) = \frac{r(x)^\beta}{\sum_y s(x, y)}$$


A total of 3 parameters:  $\alpha, \beta, \sigma$

---



---

# NEUROEVOLUTION






---

# A QUICK AND NOT-TOO-INACCURATE RECAP ON NEURAL NETWORKS

---

Input layer




Output layer

Hidden layer



# A QUICK AND NOT-TOO-INACCURATE RECAP ON NEURAL NETWORKS






---

# A QUICK AND NOT-TOO-INACCURATE RECAP ON NEURAL NETWORKS

---






---

# A QUICK AND NOT-TOO-INACCURATE RECAP ON NEURAL NETWORKS

---



All neurons in the convolution layer  
share the same weights

This is in contrast with fully connected layers, where

- 1) all connections are present
- 2) weight can be different across neurons

---



---

# NEURAL NETWORK TRAINING

---

- Training is usually performed via (stochastic) gradient descend
  - Main idea: find in which direction the weights can be modified to reduce the error by differentiating the error w.r.t. the weights
  - However, only the weights are learned
  - The architecture of the network (i.e., number, size, and type of hidden layers) and the hyperparameters are selected manually
-



---

# WHAT TO EVOLVE

---

- Weights (but this can also be done with backpropagation)
  - Activation functions
  - Hyperparameters (e.g., momentum, learning rate, dropout)
  - Architecture (e.g., connections, types of layer)
-



---

# HISTORY (CLASSIC NEUROEVOLUTION)

---

- History: anything before the advent of large/deep networks
  - Evolution of weights and topology
  - The “unit” was the single neuron/single weight
  - Currently unfeasible due to the large number of neurons/weights/layers
-



---

# NEAT

## NEUROEVOLUTION OF AUGMENTING TOPOLOGIES

---

- Think “direct encoding of a graph structure” where both the structure of the network and the weight are evolved at the same time
  - The initial population start “simple” (no hidden nodes). Evolution might add new nodes
  - Each gene has a “birth time” that tracks when it was introduced. This feature is used to allow crossover
  - Features are protected with speciation. That is, at each generation only individuals in the same specie can mate
-



---

# INDIRECT ENCODINGS

---

- One of the advantages of indirect encoding is a compact representations
  - But also the ability to express regularities, as the one present in convolution layers
  - We will see two examples without many details:  
HyperNEAT and DENSER
-




---

# CPPN

## COMPOSITIONAL PATTERN-PRODUCING NETWORKS

---




A “network” of functions that maps a position  $(i,j)$  in a 2D plane (more dimensions possible) to a value

This can be evolved as a “classical” network.  
So, where's the indirect encoding?



# HYPERNEAT





---

# DENSER

## DEEP EVOLUTIONARY NETWORK STRUCTURED REPRESENTATION

---

- A layer-based approach:  
the number, type, and parameters of the layers are evolved
  - The weights are obtained via backpropagation
  - A two-levels approach:
    - the sequence of layer and a “general” type is encoded by a GA
    - The parameters of each layer are generated by grammatical evolution (GE), in particular dynamic structured GE
  - Only the best-performing networks are trained for more than a few epochs
-